

Sleep Quality Analysis

Overview

Sleep plays a critical role in maintaining overall health, cognitive function, and emotional well-being. However, factors such as occupation, stress levels, physical activity, and lifestyle habits can significantly influence the quality and duration of sleep. With increasing awareness of sleep disorders like insomnia and sleep apnea, it is essential to analyze patterns and potential risk factors contributing to poor sleep quality.

The goal is to analyze relationships between these variables and to identify key indicators or patterns that may predict or influence sleep disorders. In this project, I will be exploring and analyzing the sleep quality data in order to identify lifestyle, health, and demographic factors that most strongly correlate with poor sleep quality. Based on the results of my analysis, I will make conclusions about which factors most strongly correlate with poor sleep quality and make recommendations for people who would like to get a better night's sleep.

The data has provided dataset of sleep and lifestyle metrics for 374 individuals. This dataset contains average values for each person.

Column Description | Column Description | ... | Person ID | An identifier for each individual. | Gender | The gender of the person | Male/Female | | Age | The age of the person in years. | Occupation | The occupation or profession of the person. | Sleep Duration (hours) | The average number of hours the person sleeps per day. | Quality of Sleep (scale: 1-10) | A subjective rating of the quality of sleep, ranging from 1 to 10. | Physical Activity Level (minutes/day) | The average number of minutes the person engages in physical activity daily. | Stress Level (scale: 1-10) | A subjective rating of the stress level experienced by the person, ranging from 1 to 10. | BMI Category | The BMI category of the person (e.g., Underweight, Normal, Overweight). | Blood Pressure (systolic/diastolic) | The average blood pressure measurement of the person, indicated as systolic pressure over diastolic pressure. | Heart Rate (bpm) | The average resting heart rate of the person in beats per minute. | Daily Steps | The average number of steps the person takes per day. | Sleep Disorder | The presence or absence of a sleep disorder in the person (None, Insomnia, Sleep Apnea). |

Importing relevant libraries

```
In [165]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
In [149]: raw_data = pd.read_csv("Desktop/sleep_health_data.csv")
raw_data
```

```
Out[149]: Person ID  Gender  Age  Occupation  Sleep Duration  Quality of Sleep  Physical Activity Level  Stress Level  BMI Category  Blood Pressure  Heart Rate  Daily Steps  Sleep Disorder
0      1   Male  27  Software Engineer    6.1          6        42       6  Overweight  126/83     77    4200      NaN
1      2   Male  28         Doctor    6.2          6        60       8    Normal  125/80     75   10000      NaN
2      3   Male  28         Doctor    6.2          6        60       8    Normal  125/80     75   10000      NaN
3      4   Male  28 Sales Representative    5.9          4        30       8    Obese  140/90     85   3000  Sleep Apnea
4      5   Male  28 Sales Representative    5.9          4        30       8    Obese  140/90     85   3000  Sleep Apnea
...      ...
369    370 Female  59           Nurse    8.1          9        75       3  Overweight  140/95     68    7000  Sleep Apnea
370    371 Female  59           Nurse    8.0          9        75       3  Overweight  140/95     68    7000  Sleep Apnea
371    372 Female  59           Nurse    8.1          9        75       3  Overweight  140/95     68    7000  Sleep Apnea
372    373 Female  59           Nurse    8.1          9        75       3  Overweight  140/95     68    7000  Sleep Apnea
373    374 Female  59           Nurse    8.1          9        75       3  Overweight  140/95     68    7000  Sleep Apnea
```

374 rows x 13 columns

Exploring the data to check missing values

```
In [166]: raw_data.isnull().sum()
percentage_missing = raw_data.isnull().mean()
percentage_missing = percentage_missing[percentage_missing > 0]
percentage_missing
```

```
Out[166]: Sleep Disorder: 0.58561
```

```
df.dropna(inplace=True)
```

```
In [170]: raw_data['Gender'].value_counts()
```

```
Out[170]: Gender
Male    189
Female  185
Name: count, dtype: int64
```

Observation: The gender distribution in the survey is approximately equal, with a balanced representation of men and women.

Healthcare workers (nurses, doctors) make up the majority of survey participants. Note that this can introduce bias into the analysis as occupations are not evenly represented. Also, since there are only 2 'Sales Representatives', I am combining them into the 'Salesperson' category.

```
In [171]: raw_data['Occupation'].value_counts()
```

```
Out[171]: Occupation
Nurse      73
Doctor     71
Engineer   63
Lawyer     47
Teacher    46
Accountant 37
Salesperson 32
Software Enginner 4
Scientist   4
Sales Representative 2
Manager    1
Name: count, dtype: int64
```

```
In [174]: raw_data['Occupation'] = raw_data['Occupation'].replace(['Sales Representative': 'Salesperson'])
```

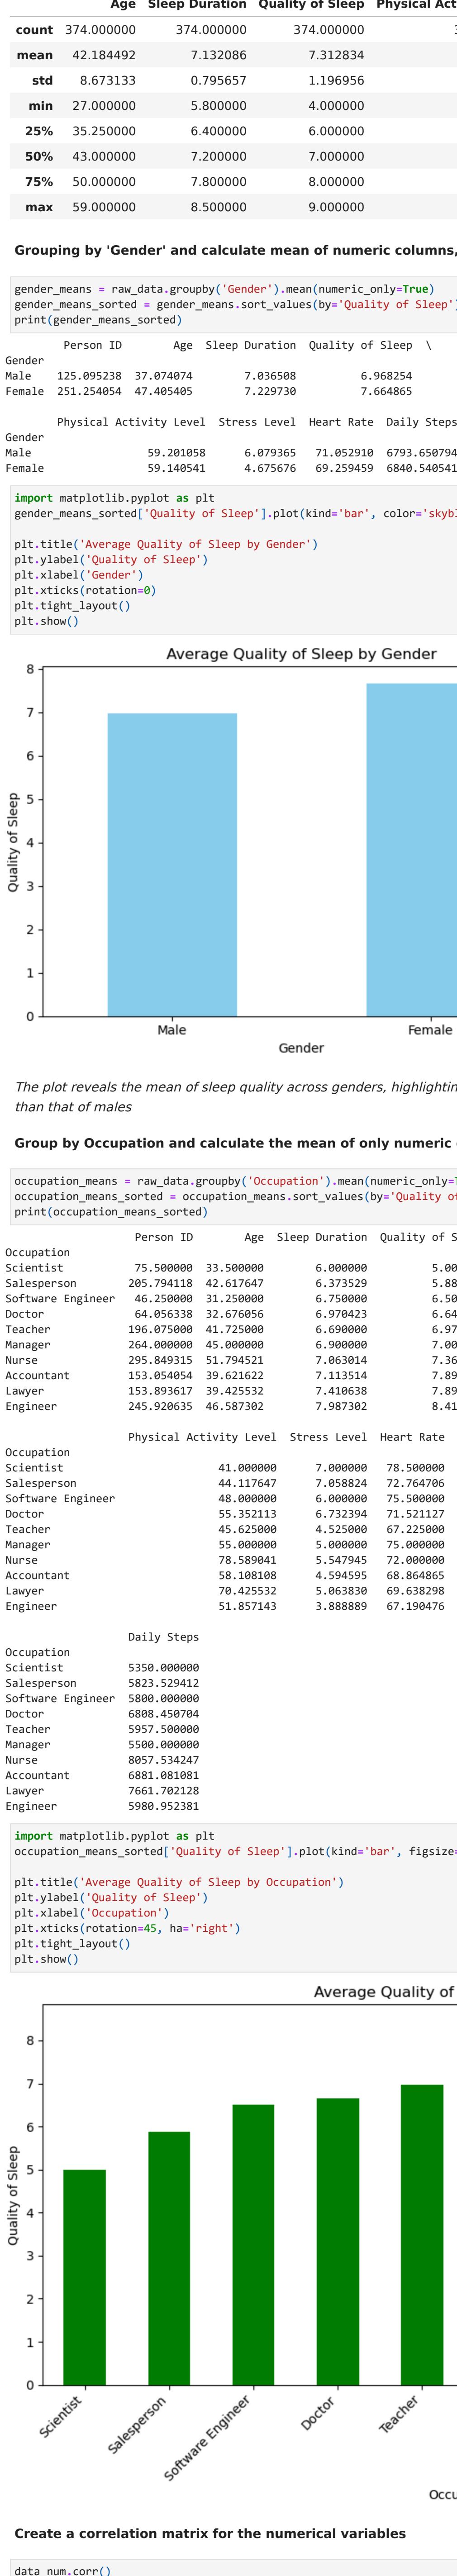
```
raw_data['Occupation'].value_counts()
```

```
Out[174]: Occupation
Nurse      73
Doctor     71
Engineer   63
Lawyer     47
Teacher    46
Accountant 37
Salesperson 32
Software Enginner 4
Scientist   4
Sales Representative 2
Manager    1
Name: count, dtype: int64
```

```
In [175]: plt.hist(raw_data.Age, color='green')
```

```
Out[175]: array([20, 26, 55, 25, 28, 14, 37, 26, 33, 1], array([27, 30, 33, 4, 36, 9, 43, 46, 42, 49, 52, 6, 55, 8, 99, 1]),
```

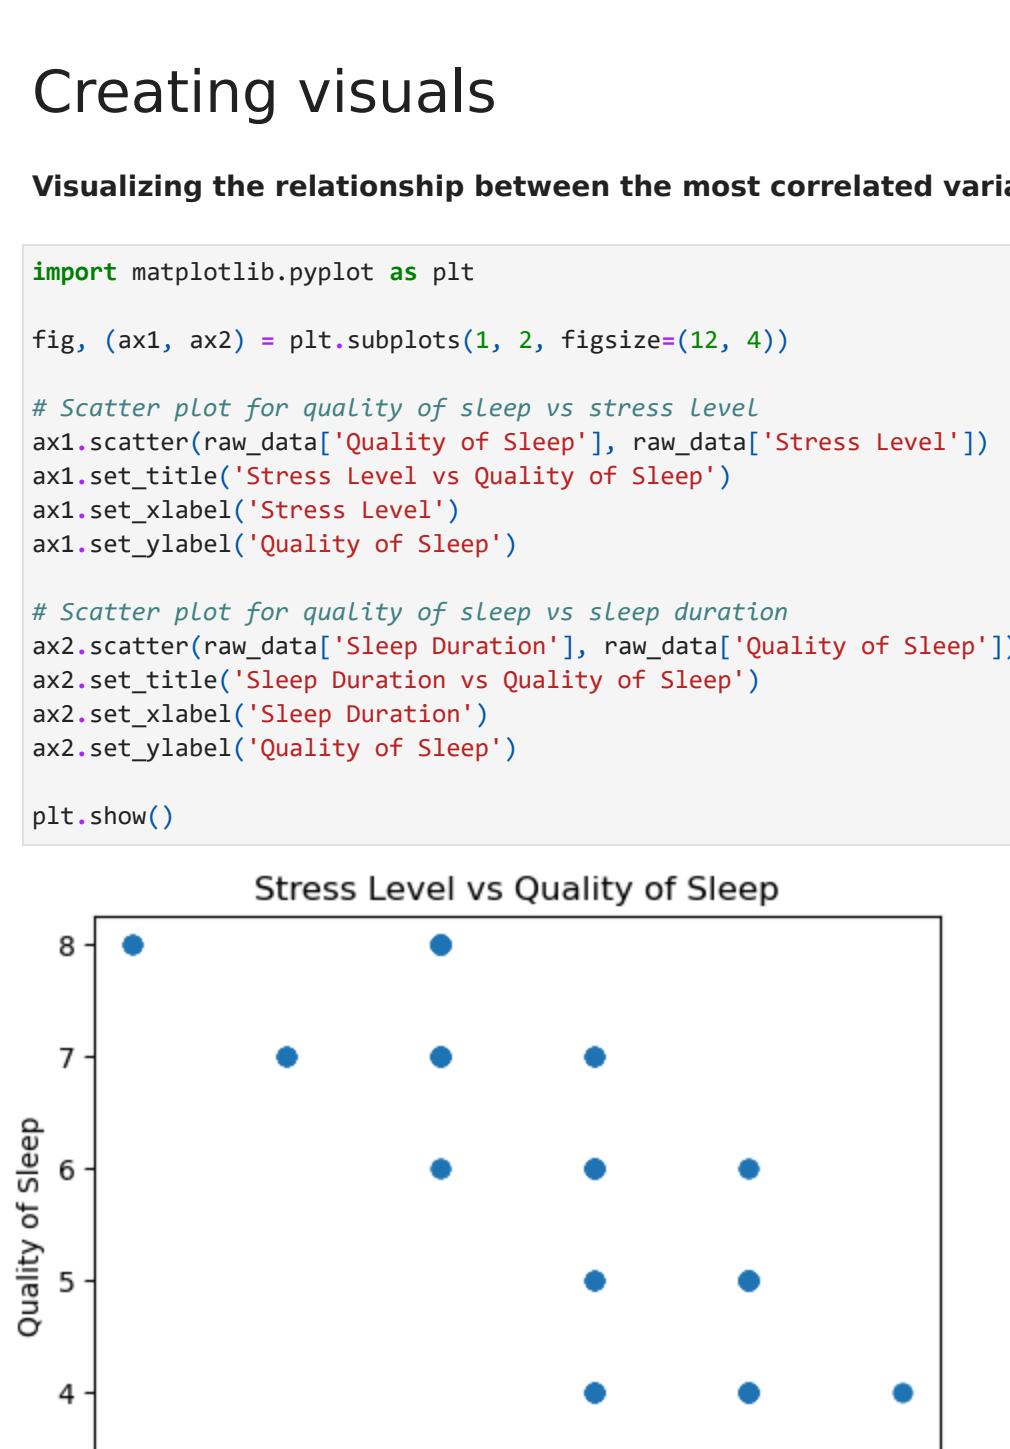
```
<chararray: object of 19 artists>
```



```
In [177]: plt.hist(raw_data['Daily Steps'], color='pink')
```

```
Out[177]: array([0, 3000, 4000, 5000, 6000, 7000, 8000, 9000, 10000], array([10, 65, 70, 75, 80, 85, 90, 95, 100]),
```

```
<chararray: object of 19 artists>
```



```
In [178]: raw_data[['Age', 'Sleep Duration', 'Quality of Sleep', 'Physical Activity Level', 'Stress Level', 'Heart Rate', 'Daily Steps']]
```

```
View basic stats of all variables in data_num
```

```
In [186]: data_num.describe()
```

```
Out[186]: Age  Sleep Duration  Quality of Sleep  Physical Activity Level  Stress Level  Heart Rate  Daily Steps
```

```
count    374.000000  374.000000    7.312034  59.111213  5.385027  374.000000  374.000000
```

```
mean    42.184492  71.732000    7.312034  59.111213  5.385027  374.000000  374.000000
```

```
std     8.673133  0.795657  1.196956  20.810804  1.774526  4.135676  1617.915679
```

```
min     27.000000  5.800000  4.000000  30.000000  65.000000  3000.000000
```

```
25%    35.250000  6.400000  6.000000  45.000000  4.000000  68.000000  5600.000000
```

```
50%    45.000000  7.200000  7.000000  60.000000  5.000000  70.000000  7000.000000
```

```
75%    50.000000  7.800000  8.000000  75.000000  7.000000  72.000000  8000.000000
```

```
max    59.000000  8.500000  9.000000  90.000000  8.000000  88.000000  10000.000000
```

```
Grouping by 'Gender' and calculate mean of numeric columns, Sort the resulting averages by 'Quality of Sleep'
```

```
In [188]: gender_mean = raw_data.groupby('Gender').mean(numeric_only=True)
```

```
gender_mean_sorted = gender_mean.sort_values(by='Quality of Sleep')
```

```
print(gender_mean_sorted)
```

```
Out[188]: Person ID  Age  Sleep Duration  Quality of Sleep  Physical Activity Level  Stress Level  Heart Rate  Daily Steps
```

```
Gender
Male    42.184492  71.732000    7.312034  59.111213  5.385027  374.000000  374.000000
```

```
Female   37.047074  70.457076  7.083213  58.832130  5.192890  374.000000  374.000000
```

```
Normal   42.184492  71.732000    7.312034  59.111213  5.385027  374.000000  374.000000
```

```
Overweight 37.047074  70.457076  6.979423  58.832130  5.192890  374.000000  374.000000
```

```
Obese    37.047074  70.457076  6.647887  58.832130  5.192890  374.000000  374.000000
```

```
Name: count, dtype: int64
```

```
Combine 'Normal' and 'Normal Weight' categories since they appear to mean the same thing
```

```
In [189]: raw_data['BMI Category'].value_counts()
```

```
Out[189]: BMI Category
Normal    216
Overweight 144
Underweight 104
Obese     10
Name: count, dtype: int64
```

```
Based on initial exploration, the dataset appears clean and free of any significant outliers. I will next look at some basic stats and create groupings to examine relationships between variables.
```

```
Summarizing the data by creating a datafarme of just the numerical variables
```

```
In [190]: data_num = raw_data[['Age', 'Sleep Duration', 'Quality of Sleep', 'Physical Activity Level', 'Stress Level', 'Heart Rate', 'Daily Steps']]
```

```
View basic stats of all variables in data_num
```

```
In [186]: data_num.describe()
```

```
Out[186]: Age  Sleep Duration  Quality of Sleep  Physical Activity Level  Stress Level  Heart Rate  Daily Steps
```

```
count    374.000000  374.000000    7.312034  59.111213  5.385027  374.000000  374.000000
```

```
mean    42.184492  71.732000    7.312034  59.111213  5.385027  374.000000  374.000000
```

```
std     8.673133  0.795657  1.196956  20.810804  1.774526  4.135676  1617.915679
```

```
min     27.000000  5.800000  4.000000  30.000000  65.000000  3000.000000
```

```
25%    35.250000  6.400000  6.000000  45.000000  4.000000  68.000000  5600.000000
```

```
50%    45.000000  7.200000  7.000000  60.000000  5.000000  70.000000  7000.000000
```

```
75%    50.000000  7.800000  8.000000  75.000000  7.000000  72.000000  8000.000000
```

```
max    59.000000  8.500000  9.000000  90.000000  8.000000  88.000000  10000.000000
```

```
Grouping by 'Gender' and calculate mean of numeric columns, Sort the resulting averages by 'Quality of Sleep'
```

```
In [188]: gender_mean = raw_data.groupby('Gender').mean(numeric_only=True)
```

```
gender_mean_sorted = gender_mean.sort_values(by='Quality of Sleep')
```

```
print(gender_mean_sorted)
```

```
Out[188]: Person ID  Age  Sleep Duration  Quality of Sleep  Physical Activity Level  Stress Level  Heart Rate  Daily Steps
```

```
Gender
Male    42.184492  71.732000    7.312034  59.111213  5.385027  374.000000  374.000000
```

```
Female   37.047074  70.457076  6.883213  58.832130  5.192890  374.000000  374.000000
```

```
Normal   42.184492  71.732000    7.312034  59.111213  5.385027  374.000000  374.000000
```

```
Overweight 37.047074  70.457076  6.647887  58.832130  5.192890  374.000000  374.000000
```

```
Obese    37.047074  70.457076  6.423344  58.832130  5.192890  374.000000  374.000000
```

```
Name: count, dtype: int64
```

```
Combine 'Normal' and 'Normal Weight' categories since they appear to mean the same thing
```

```
In [189]: raw_data['BMI Category'].value_counts()
```

```
Out[189]: BMI Category
Normal    216
Overweight 144
Underweight 104
Obese     10
Name: count, dtype: int64
```

```
Based on initial exploration, the dataset appears clean and free of any significant outliers. I will next look at some basic stats and create groupings to examine relationships between variables.
```

```
Summarizing the data by creating a datafarme of just the numerical variables
```

```
In [190]: data_num = raw_data[['Age', 'Sleep Duration', 'Quality of Sleep', 'Physical Activity Level', 'Stress Level', 'Heart Rate', 'Daily Steps']]
```

```
View basic stats of all variables in data_num
```

```
In [186]: data_num.describe()
```

```
Out[186]: Age  Sleep Duration  Quality of Sleep  Physical Activity Level  Stress Level  Heart Rate  Daily Steps
```

```
count    374.000000  374.000000    7.312034  59.111213  5.385027  374.000000  374.000000
```

```
mean    42.184492  71.732000    7.312034  59.111213  5.385027  374.00
```